

In [1]:

```

1 # Importing Pandas and NumPy
2 import pandas as pd
3 import numpy as np

```

In [2]:

```

1 # Importing dataset
2 data = pd.read_csv('D:/Python/Dataset/nrippner-titanic-disaster-dataset/titanic.csv')

```

In [3]:

```

1 data

```

Out[3]:

	pclass	survived	name	sex	age	sibsp	parch	ticket	fare	cabin	e
0	1.0	1.0	Allen, Miss. Elisabeth Walton	female	29.0000	0.0	0.0	24160	211.3375	B5	
1	1.0	1.0	Allison, Master. Hudson Trevor	male	0.9167	1.0	2.0	113781	151.5500	C22 C26	
2	1.0	0.0	Allison, Miss. Helen Loraine	female	2.0000	1.0	2.0	113781	151.5500	C22 C26	
3	1.0	0.0	Allison, Mr. Hudson Joshua Creighton	male	30.0000	1.0	2.0	113781	151.5500	C22 C26	
4	1.0	0.0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	female	25.0000	1.0	2.0	113781	151.5500	C22 C26	
...
1305	3.0	0.0	Zabour, Miss. Thamine	female	NaN	1.0	0.0	2665	14.4542	NaN	
1306	3.0	0.0	Zakarian, Mr. Mapriededer	male	26.5000	0.0	0.0	2656	7.2250	NaN	
1307	3.0	0.0	Zakarian, Mr. Ortin	male	27.0000	0.0	0.0	2670	7.2250	NaN	
1308	3.0	0.0	Zimmerman, Mr. Leo	male	29.0000	0.0	0.0	315082	7.8750	NaN	
1309	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

1310 rows × 14 columns



In [4]:

```
1 data.describe()
```

Out[4]:

	pclass	survived	age	sibsp	parch	fare	bo
count	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000	1308.000000	121.0000
mean	2.294882	0.381971	29.881135	0.498854	0.385027	33.295479	160.8099
std	0.837836	0.486055	14.413500	1.041658	0.865560	51.758668	97.6969
min	1.000000	0.000000	0.166700	0.000000	0.000000	0.000000	1.0000
25%	2.000000	0.000000	21.000000	0.000000	0.000000	7.895800	72.0000
50%	3.000000	0.000000	28.000000	0.000000	0.000000	14.454200	155.0000
75%	3.000000	1.000000	39.000000	1.000000	0.000000	31.275000	256.0000
max	3.000000	1.000000	80.000000	8.000000	9.000000	512.329200	328.0000

In [5]:

```
1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1310 entries, 0 to 1309
Data columns (total 14 columns):
pclass      1309 non-null float64
survived     1309 non-null float64
name        1309 non-null object
sex         1309 non-null object
age         1046 non-null float64
sibsp       1309 non-null float64
parch       1309 non-null float64
ticket      1309 non-null object
fare        1308 non-null float64
cabin       295 non-null object
embarked    1307 non-null object
boat        486 non-null object
body        121 non-null float64
home.dest   745 non-null object
dtypes: float64(7), object(7)
memory usage: 143.4+ KB
```

In [6]:

```
1 data['sex'].value_counts()
```

Out[6]:

```
male      843
female    466
Name: sex, dtype: int64
```

In [7]:

```

1 # Converting Male to 1 and Female to 0
2 data['sex'] = data['sex'].map({'male': 1, 'female': 0})
3 #The variable was imported as a string we need to convert it to float
4 #data['sex'] =pd.to_numeric(data['sex'],errors='coerce')

```

In [8]:

```

1 data

```

Out[8]:

	pclass	survived	name	sex	age	sibsp	parch	ticket	fare	cabin	em
0	1.0	1.0	Allen, Miss. Elisabeth Walton	0.0	29.0000	0.0	0.0	24160	211.3375	B5	
1	1.0	1.0	Allison, Master. Hudson Trevor	1.0	0.9167	1.0	2.0	113781	151.5500	C22 C26	
2	1.0	0.0	Allison, Miss. Helen Loraine	0.0	2.0000	1.0	2.0	113781	151.5500	C22 C26	
3	1.0	0.0	Allison, Mr. Hudson Joshua Creighton	1.0	30.0000	1.0	2.0	113781	151.5500	C22 C26	
4	1.0	0.0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	0.0	25.0000	1.0	2.0	113781	151.5500	C22 C26	
...
1305	3.0	0.0	Zabour, Miss. Thamine	0.0	NaN	1.0	0.0	2665	14.4542	NaN	
1306	3.0	0.0	Zakarian, Mr. Mapriededer	1.0	26.5000	0.0	0.0	2656	7.2250	NaN	
1307	3.0	0.0	Zakarian, Mr. Ortin	1.0	27.0000	0.0	0.0	2670	7.2250	NaN	
1308	3.0	0.0	Zimmerman, Mr. Leo	1.0	29.0000	0.0	0.0	315082	7.8750	NaN	
1309	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

1310 rows × 14 columns



In [9]:

```
1 data.describe()
```

Out[9]:

	pclass	survived	sex	age	sibsp	parch	1
count	1309.000000	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000	1308.000
mean	2.294882	0.381971	0.644003	29.881135	0.498854	0.385027	33.295
std	0.837836	0.486055	0.478997	14.413500	1.041658	0.865560	51.758
min	1.000000	0.000000	0.000000	0.166700	0.000000	0.000000	0.000
25%	2.000000	0.000000	0.000000	21.000000	0.000000	0.000000	7.895
50%	3.000000	0.000000	1.000000	28.000000	0.000000	0.000000	14.454
75%	3.000000	1.000000	1.000000	39.000000	1.000000	0.000000	31.275
max	3.000000	1.000000	1.000000	80.000000	8.000000	9.000000	512.329

In [10]:

```
1 # Creating a dummy variable for the variable 'Contract' and dropping the first one.
2 cont = pd.get_dummies(data['embarked'],prefix='embarked',drop_first=True)
3 #Adding the results to the master dataframe
4 data = pd.concat([data,cont],axis=1)
```

In [11]:

```
1 data
```

Out[11]:

	pclass	survived	name	sex	age	sibsp	parch	ticket	fare	cabin	em
0	1.0	1.0	Allen, Miss. Elisabeth Walton	0.0	29.0000	0.0	0.0	24160	211.3375	B5	
1	1.0	1.0	Allison, Master. Hudson Trevor	1.0	0.9167	1.0	2.0	113781	151.5500	C22 C26	
2	1.0	0.0	Allison, Miss. Helen Loraine	0.0	2.0000	1.0	2.0	113781	151.5500	C22 C26	
3	1.0	0.0	Allison, Mr. Hudson Joshua Creighton	1.0	30.0000	1.0	2.0	113781	151.5500	C22 C26	
4	1.0	0.0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	0.0	25.0000	1.0	2.0	113781	151.5500	C22 C26	
...
1305	3.0	0.0	Zabour, Miss. Thamine	0.0	NaN	1.0	0.0	2665	14.4542	NaN	
1306	3.0	0.0	Zakarian, Mr. Mapriededer	1.0	26.5000	0.0	0.0	2656	7.2250	NaN	
1307	3.0	0.0	Zakarian, Mr. Ortin	1.0	27.0000	0.0	0.0	2670	7.2250	NaN	
1308	3.0	0.0	Zimmerman, Mr. Leo	1.0	29.0000	0.0	0.0	315082	7.8750	NaN	
1309	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

1310 rows × 16 columns



In [12]:

```
1 data.describe()
```

Out[12]:

	pclass	survived	sex	age	sibsp	parch	1
count	1309.000000	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000	1308.000
mean	2.294882	0.381971	0.644003	29.881135	0.498854	0.385027	33.295
std	0.837836	0.486055	0.478997	14.413500	1.041658	0.865560	51.758
min	1.000000	0.000000	0.000000	0.166700	0.000000	0.000000	0.000
25%	2.000000	0.000000	0.000000	21.000000	0.000000	0.000000	7.895
50%	3.000000	0.000000	1.000000	28.000000	0.000000	0.000000	14.454
75%	3.000000	1.000000	1.000000	39.000000	1.000000	0.000000	31.275
max	3.000000	1.000000	1.000000	80.000000	8.000000	9.000000	512.329

In [13]:

```
1 data['ticket'].value_counts()
```

Out[13]:

```
CA. 2343      11
1601          8
CA 2144        8
347077         7
S.O.C. 14879   7
..
364850         1
315083         1
350035         1
350404         1
31028          1
Name: ticket, Length: 929, dtype: int64
```

In [14]:

```
1 # We have created dummies for the below variables, so we can drop them
2 data = data.drop(['embarked'], 1)
```

In [15]:

```
1 data
```

Out[15]:

	pclass	survived	name	sex	age	sibsp	parch	ticket	fare	cabin	boat
0	1.0	1.0	Allen, Miss. Elisabeth Walton	0.0	29.0000	0.0	0.0	24160	211.3375	B5	
1	1.0	1.0	Allison, Master. Hudson Trevor	1.0	0.9167	1.0	2.0	113781	151.5500	C22 C26	1
2	1.0	0.0	Allison, Miss. Helen Loraine	0.0	2.0000	1.0	2.0	113781	151.5500	C22 C26	Na
3	1.0	0.0	Allison, Mr. Hudson Joshua Creighton	1.0	30.0000	1.0	2.0	113781	151.5500	C22 C26	Na
4	1.0	0.0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	0.0	25.0000	1.0	2.0	113781	151.5500	C22 C26	Na
...
1305	3.0	0.0	Zabour, Miss. Thamine	0.0	NaN	1.0	0.0	2665	14.4542	NaN	Na
1306	3.0	0.0	Zakarian, Mr. Mapriededer	1.0	26.5000	0.0	0.0	2656	7.2250	NaN	Na
1307	3.0	0.0	Zakarian, Mr. Ortin	1.0	27.0000	0.0	0.0	2670	7.2250	NaN	Na
1308	3.0	0.0	Zimmerman, Mr. Leo	1.0	29.0000	0.0	0.0	315082	7.8750	NaN	Na
1309	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	Na

1310 rows × 15 columns



In [16]:

```

1 # Checking outliers at 25%,50%,75%,90%,95% and 99%
2 data.describe(percentiles=[.25,.5,.75,.90,.95,.99])

```

Out[16]:

	pclass	survived	sex	age	sibsp	parch	1
count	1309.000000	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000	1308.000
mean	2.294882	0.381971	0.644003	29.881135	0.498854	0.385027	33.295
std	0.837836	0.486055	0.478997	14.413500	1.041658	0.865560	51.758
min	1.000000	0.000000	0.000000	0.166700	0.000000	0.000000	0.000
25%	2.000000	0.000000	0.000000	21.000000	0.000000	0.000000	7.895
50%	3.000000	0.000000	1.000000	28.000000	0.000000	0.000000	14.454
75%	3.000000	1.000000	1.000000	39.000000	1.000000	0.000000	31.275
90%	3.000000	1.000000	1.000000	50.000000	1.000000	2.000000	78.050
95%	3.000000	1.000000	1.000000	57.000000	2.000000	2.000000	133.650
99%	3.000000	1.000000	1.000000	65.000000	5.000000	4.000000	262.375
max	3.000000	1.000000	1.000000	80.000000	8.000000	9.000000	512.329

In [17]:

```

1 # Adding up the missing values (column-wise)
2 data.isnull().sum()

```

Out[17]:

```

pclass      1
survived     1
name         1
sex          1
age        264
sibsp        1
parch        1
ticket       1
fare         2
cabin     1015
boat        824
body       1189
home.dest    565
embarked_Q     0
embarked_S     0
dtype: int64

```


In [18]:

```
1 # Checking the percentage of missing values
2 round(100*(data.isnull().sum()/len(data.index)), 2)
```

Out[18]:

pclass	0.08
survived	0.08
name	0.08
sex	0.08
age	20.15
sibsp	0.08
parch	0.08
ticket	0.08
fare	0.15
cabin	77.48
boat	62.90
body	90.76
home.dest	43.13
embarked_Q	0.00
embarked_S	0.00

dtype: float64

In [88]:

```
1 data['home.dest'].value_counts()
```

```
-----
KeyError                                Traceback (most recent call last)
~\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self, key, method, tolerance)
    2896         try:
-> 2897             return self._engine.get_loc(key)
    2898         except KeyError:
```

```
pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()
```

```
pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()
```

```
pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()
```

```
pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()
```

```
KeyError: 'home.dest'
```

During handling of the above exception, another exception occurred:

```
KeyError                                Traceback (most recent call last)
<ipython-input-88-114bb5a0bd52> in <module>
----> 1 data['home.dest'].value_counts()

~\Anaconda3\lib\site-packages\pandas\core\frame.py in __getitem__(self, key)
    2978         if self.columns.nlevels > 1:
    2979             return self._getitem_multilevel(key)
-> 2980         indexer = self.columns.get_loc(key)
    2981         if is_integer(indexer):
    2982             indexer = [indexer]
```

```
~\Anaconda3\lib\site-packages\pandas\core\indexes\base.py in get_loc(self, key, method, tolerance)
    2897         return self._engine.get_loc(key)
    2898         except KeyError:
-> 2899         return self._engine.get_loc(self._maybe_cast_indexer(key))
    2900         indexer = self.get_indexer([key], method=method, tolerance=tolerance)
    2901         if indexer.ndim > 1 or indexer.size > 1:
```

```
pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()
```

```
pandas\_libs\index.pyx in pandas._libs.index.IndexEngine.get_loc()
```

```
pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()
```

```
pandas\_libs\hashtable_class_helper.pxi in pandas._libs.hashtable.PyObjectHashTable.get_item()
```

```
KeyError: 'home.dest'
```

In [20]:

```
1 # We have some unwanted columns , so we can drop them
2 data = data.drop(['home.dest','cabin','boat'], 1)
```

In [21]:

```
1 data
```

Out[21]:

	pclass	survived	name	sex	age	sibsp	parch	ticket	fare	body	eml
0	1.0	1.0	Allen, Miss. Elisabeth Walton	0.0	29.0000	0.0	0.0	24160	211.3375	NaN	
1	1.0	1.0	Allison, Master. Hudson Trevor	1.0	0.9167	1.0	2.0	113781	151.5500	NaN	
2	1.0	0.0	Allison, Miss. Helen Loraine	0.0	2.0000	1.0	2.0	113781	151.5500	NaN	
3	1.0	0.0	Allison, Mr. Hudson Joshua Creighton	1.0	30.0000	1.0	2.0	113781	151.5500	135.0	
4	1.0	0.0	Allison, Mrs. Hudson J C (Bessie Waldo Daniels)	0.0	25.0000	1.0	2.0	113781	151.5500	NaN	
...
1305	3.0	0.0	Zabour, Miss. Thamine	0.0	NaN	1.0	0.0	2665	14.4542	NaN	
1306	3.0	0.0	Zakarian, Mr. Mapriededer	1.0	26.5000	0.0	0.0	2656	7.2250	304.0	
1307	3.0	0.0	Zakarian, Mr. Ortin	1.0	27.0000	0.0	0.0	2670	7.2250	NaN	
1308	3.0	0.0	Zimmerman, Mr. Leo	1.0	29.0000	0.0	0.0	315082	7.8750	NaN	
1309	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

1310 rows × 12 columns



In [22]:

```
1 data['age'].tail(50)
```

Out[22]:

```
1260    18.0
1261    63.0
1262     NaN
1263    11.5
1264    40.5
1265    10.0
1266    36.0
1267    30.0
1268     NaN
1269    33.0
1270    28.0
1271    28.0
1272    47.0
1273    18.0
1274    31.0
1275    16.0
1276    31.0
1277    22.0
1278    20.0
1279    14.0
1280    22.0
1281    22.0
1282     NaN
1283     NaN
1284     NaN
1285    32.5
1286    38.0
1287    51.0
1288    18.0
1289    21.0
1290    47.0
1291     NaN
1292     NaN
1293     NaN
1294    28.5
1295    21.0
1296    27.0
1297     NaN
1298    36.0
1299    27.0
1300    15.0
1301    45.5
1302     NaN
1303     NaN
1304    14.5
1305     NaN
1306    26.5
1307    27.0
1308    29.0
1309     NaN
```

```
Name: age, dtype: float64
```

In [23]:

```
1 data.describe()
```

Out[23]:

	pclass	survived	sex	age	sibsp	parch	1
count	1309.000000	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000	1308.000
mean	2.294882	0.381971	0.644003	29.881135	0.498854	0.385027	33.295
std	0.837836	0.486055	0.478997	14.413500	1.041658	0.865560	51.758
min	1.000000	0.000000	0.000000	0.166700	0.000000	0.000000	0.000
25%	2.000000	0.000000	0.000000	21.000000	0.000000	0.000000	7.895
50%	3.000000	0.000000	1.000000	28.000000	0.000000	0.000000	14.454
75%	3.000000	1.000000	1.000000	39.000000	1.000000	0.000000	31.275
max	3.000000	1.000000	1.000000	80.000000	8.000000	9.000000	512.329

In [24]:

```
1 data.describe(percentiles=[.25,.5,.75,.90,.95,.99])
```

Out[24]:

	pclass	survived	sex	age	sibsp	parch	1
count	1309.000000	1309.000000	1309.000000	1046.000000	1309.000000	1309.000000	1308.000
mean	2.294882	0.381971	0.644003	29.881135	0.498854	0.385027	33.295
std	0.837836	0.486055	0.478997	14.413500	1.041658	0.865560	51.758
min	1.000000	0.000000	0.000000	0.166700	0.000000	0.000000	0.000
25%	2.000000	0.000000	0.000000	21.000000	0.000000	0.000000	7.895
50%	3.000000	0.000000	1.000000	28.000000	0.000000	0.000000	14.454
75%	3.000000	1.000000	1.000000	39.000000	1.000000	0.000000	31.275
90%	3.000000	1.000000	1.000000	50.000000	1.000000	2.000000	78.050
95%	3.000000	1.000000	1.000000	57.000000	2.000000	2.000000	133.650
99%	3.000000	1.000000	1.000000	65.000000	5.000000	4.000000	262.375
max	3.000000	1.000000	1.000000	80.000000	8.000000	9.000000	512.329

In [25]:

```
1 data.isnull().sum()
```

Out[25]:

```
pclass      1
survived     1
name         1
sex          1
age         264
sibsp        1
parch        1
ticket       1
fare         2
body        1189
embarked_Q   0
embarked_S   0
dtype: int64
```

In [26]:

```
1 data['age'].fillna(data['age'].mean(),inplace=True)
```

In [27]:

```
1 data['fare'].fillna(data['fare'].mean(),inplace=True)
```

In [28]:

```
1 data.isnull().sum()
```

Out[28]:

```
pclass      1
survived     1
name         1
sex          1
age          0
sibsp        1
parch        1
ticket       1
fare         0
body        1189
embarked_Q   0
embarked_S   0
dtype: int64
```

In [29]:

```
1 # Normalising continuous features
2 df = data[['sex','age','fare']]
```

In [30]:

```
1 normalized_df=(df-df.mean())/df.std()
```

In [31]:

```
1 data = data.drop(['sex', 'age', 'fare'], 1)
```

In [32]:

```
1 data = pd.concat([data, normalized_df], axis=1)
```

In [33]:

```
1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1310 entries, 0 to 1309
Data columns (total 12 columns):
pclass      1309 non-null float64
survived     1309 non-null float64
name        1309 non-null object
sibsp       1309 non-null float64
parch       1309 non-null float64
ticket      1309 non-null object
body        121 non-null float64
embarked_Q  1310 non-null uint8
embarked_S  1310 non-null uint8
sex         1309 non-null float64
age         1310 non-null float64
fare        1310 non-null float64
dtypes: float64(8), object(2), uint8(2)
memory usage: 105.0+ KB
```

In [34]:

```
1 # Removing NaN in Survived rows
2 data = data[~np.isnan(data['survived'])]
```

In [35]:

```
1 survived = (sum(data['survived'])/len(data['survived'].index))*100
```

In [36]:

```
1 survived
```

Out[36]:

```
38.19709702062643
```

In [37]:

```
1 from sklearn.model_selection import train_test_split
```

In [38]:

```
1 # Putting feature variable to X
2 X = data.drop(['name', 'survived', 'ticket', 'body'], axis=1)
3
4 # Putting response variable to y
5 y = data['survived']
```

In [39]:

```
1 X.isnull().sum()
```

Out[39]:

```
pclass      0
sibsp       0
parch       0
embarked_Q   0
embarked_S   0
sex         0
age         0
fare        0
dtype: int64
```

In [40]:

```
1 # Splitting the data into train and test
2 X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, test_size=0.3, random_state=42)
```

In [41]:

```
1 import statsmodels.api as sm
```

In [42]:

```
1 # Let's run the model using the selected variables
2 from sklearn.linear_model import LogisticRegression
3 from sklearn import metrics
4 logsk = LogisticRegression()
5 logsk.fit(X_train, y_train)
```

C:\Users\Vicky\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
FutureWarning)

Out[42]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='warn', n_jobs=None, penalty='l2',
                    random_state=None, solver='warn', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [43]:

```
1 # Predicted probabilities
2 y_pred = logsk.predict_proba(X_test)
```


In [44]:

```
1 # Predicted probabilities
2 y_pred1 = logsk.predict(X_test)
```

In [45]:

```
1 from sklearn import metrics
2 metrics.accuracy_score( y_test, y_pred1)*100
```

Out[45]:

76.33587786259542

In [46]:

```
1 logsk1 = LogisticRegression(C=10000,penalty='l1')
2 logsk1.fit(X_train, y_train)
```

C:\Users\Vicky\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Speci
fy a solver to silence this warning.
FutureWarning)

Out[46]:

```
LogisticRegression(C=10000, class_weight=None, dual=False, fit_intercept=Tru  
e,  
                    intercept_scaling=1, l1_ratio=None, max_iter=100,  
                    multi_class='warn', n_jobs=None, penalty='l1',  
                    random_state=None, solver='warn', tol=0.0001, verbose=0,  
                    warm_start=False)
```

In [47]:

```
1 # Predicted probabilities
2 y_pred1 = logsk1.predict_proba(X_test)
```

In [48]:

1	y_pred1
---	---------

Out[48]:

```
array([[0.39286987, 0.60713013],
       [0.92442351, 0.07557649],
       [0.91191923, 0.08808077],
       [0.39401712, 0.60598288],
       [0.03604694, 0.96395306],
       [0.78627565, 0.21372435],
       [0.85648477, 0.14351523],
       [0.49035479, 0.50964521],
       [0.88571314, 0.11428686],
       [0.51616802, 0.48383198],
       [0.09139541, 0.90860459],
       [0.4032952 , 0.5967048 ],
       [0.84052779, 0.15947221],
       [0.55398624, 0.44601376],
       [0.45654732, 0.54345268],
       [0.18422854, 0.81577146],
       [0.91236422, 0.08763578],
       [0.96045032, 0.03954968].
```

In [49]:

1	metrics.accuracy_score(y_test, y_pred1)*100
---	--

ValueError

Traceback (most recent call last)

<ipython-input-49-eddbce20593e> in <module>

----> 1 metrics.accuracy_score(y_test, y_pred1)*100

~\Anaconda3\lib\site-packages\sklearn\metrics\classification.py in accuracy_score(y_true, y_pred, normalize, sample_weight)

174

175 # Compute accuracy for each possible representation

--> 176 y_type, y_true, y_pred = _check_targets(y_true, y_pred)

177 check_consistent_length(y_true, y_pred, sample_weight)

178 if y_type.startswith('multilabel'):

~\Anaconda3\lib\site-packages\sklearn\metrics\classification.py in _check_targets(y_true, y_pred)

79 if len(y_type) > 1:

80 raise ValueError("Classification metrics can't handle a mix

of {0} "

---> 81 "and {1} targets".format(type_true, type_pred))

82

83 # We can't have more than one value on y_type => The set is no more needed

ValueError: Classification metrics can't handle a mix of binary and continuous-multioutput targets

In [50]:

1	y_pred2 = logsk.predict(X_test)
---	---------------------------------

In [51]:

```
1 metrics.accuracy_score( y_test, y_pred2)*100
```

Out[51]:

76.33587786259542

In [52]:

```
1 logsk2 = LogisticRegression(C=100,penalty='elasticnet',solver='saga',l1_ratio=0)
2 logsk2.fit(X_train, y_train)
```

Out[52]:

```
LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=0, max_iter=100,
                    multi_class='warn', n_jobs=None, penalty='elasticnet',
                    random_state=None, solver='saga', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [53]:

```
1 y_pred3 = logsk.predict(X_test)
```

In [54]:

```
1 metrics.accuracy_score( y_test, y_pred3)*100
```

Out[54]:

76.33587786259542

In [55]:

```
1 logsk3 = LogisticRegression(C=10,penalty='elasticnet',solver='saga',l1_ratio=1)
2 logsk3.fit(X_train, y_train)
```

Out[55]:

```
LogisticRegression(C=10, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=1, max_iter=100,
                    multi_class='warn', n_jobs=None, penalty='elasticnet',
                    random_state=None, solver='saga', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [56]:

```
1 y_pred4 = logsk.predict(X_test)
```

In [57]:

```
1 metrics.accuracy_score( y_test, y_pred4)*100
```

Out[57]:

76.33587786259542

In [59]:

```
1 # Converting y_pred to a dataframe which is an array
2 y_pred_df = pd.DataFrame(y_pred1)
```

In [60]:

```
1 # Converting to column dataframe
2 y_pred_1 = y_pred_df.iloc[:,[1]]
```

In [61]:

```
1 # Let's see the head
2 y_pred_1.head()
```

Out[61]:

	1
0	0.607130
1	0.075576
2	0.088081
3	0.605983
4	0.963953

In [62]:

```
1 # Converting y_test to dataframe
2 y_test_df = pd.DataFrame(y_test)
```

In [63]:

```
1 # Putting CustID to index
2 y_test_df['name'] = y_test_df.index
```

In [64]:

```
1 y_test_df
```

Out[64]:

	survived	name
173	0.0	173
843	0.0	843
996	0.0	996
992	0.0	992
12	1.0	12
...
1191	0.0	1191
165	1.0	165
588	1.0	588
270	1.0	270
61	1.0	61

393 rows × 2 columns

In [65]:

```
1 # Removing index for both dataframes to append them side by side
2 y_pred_1.reset_index(drop=True, inplace=True)
3 y_test_df.reset_index(drop=True, inplace=True)
```

In [66]:

```
1 # Appending y_test_df and y_pred_1
2 y_pred_final = pd.concat([y_test_df, y_pred_1], axis=1)
```

In [67]:

```
1 # Renaming the column
2 y_pred_final = y_pred_final.rename(columns={ 1 : 'Survived_Prob'})
```

In [68]:

```
1 # Let's see the head of y_pred_final
2 y_pred_final.head()
```

Out[68]:

	survived	name	Survived_Prob
0	0.0	173	0.607130
1	0.0	843	0.075576
2	0.0	996	0.088081
3	0.0	992	0.605983
4	1.0	12	0.963953

In [69]:

```
1 # Creating new column 'predicted' with 1 if Churn_Prob>0.5 else 0
2 y_pred_final['predicted'] = y_pred_final.Survived_Prob.map( lambda x: 1 if x > 0.5 else 0)
```

In [70]:

```
1 # Let's see the head
2 y_pred_final.head()
```

Out[70]:

	survived	name	Survived_Prob	predicted
0	0.0	173	0.607130	1
1	0.0	843	0.075576	0
2	0.0	996	0.088081	0
3	0.0	992	0.605983	1
4	1.0	12	0.963953	1

In [71]:

```
1 from sklearn import metrics
```

In [74]:

```
1 # Confusion matrix
2 confusion = metrics.confusion_matrix( y_pred_final.survived, y_pred_final.predicted )
3 confusion
```

Out[74]:

```
array([[208, 45],
       [ 46, 94]], dtype=int64)
```

In [75]:

```
1 #Let's check the overall accuracy.  
2 metrics.accuracy_score( y_pred_final.survived, y_pred_final.predicted)
```

Out[75]:

0.7684478371501272

In [76]:

```
1 TP = confusion[0,0] # true positive  
2 TN = confusion[1,1] # true negatives  
3 FP = confusion[0,1] # false positives  
4 FN = confusion[1,0] # false negatives
```

In [77]:

```
1 # Let's see the sensitivity of our logistic regression model  
2 TP / float(TP+FN)
```

Out[77]:

0.8188976377952756

In [78]:

```
1 # Let us calculate specificity  
2 TN / float(TN+FP)
```

Out[78]:

0.6762589928057554

In [80]:

```
1 # Calculate false positive rate - predicting survived when customer does not survived  
2 print(FP/ float(TN+FP))
```

0.3237410071942446

In [81]:

```
1 # positive predictive value  
2 print (TP / float(TP+FP))
```

0.8221343873517787

In [82]:

```
1 # Negative predictive value  
2 print (TN / float(TN+ FN))
```

0.6714285714285714

In [85]:

```

1 def draw_roc( actual, probs ):
2     fpr, tpr, thresholds = metrics.roc_curve( actual, probs,
3                                               drop_intermediate = False )
4     auc_score = metrics.roc_auc_score( actual, probs )
5     plt.figure(figsize=(6, 4))
6     plt.plot( fpr, tpr, label='ROC curve (area = %0.2f)' % auc_score )
7     plt.plot([0, 1], [0, 1], 'k--')
8     plt.xlim([0.0, 1.0])
9     plt.ylim([0.0, 1.05])
10    plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
11    plt.ylabel('True Positive Rate')
12    plt.title('Receiver operating characteristic example')
13    plt.legend(loc="lower right")
14    plt.show()
15
16    return fpr, tpr, thresholds


```

In [87]:

```

1 # Importing matplotlib and seaborn
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 %matplotlib inline
5 draw_roc(y_pred_final.survived, y_pred_final.predicted)

```



Out[87]:

```

(array([0.          , 0.17786561, 1.          ]),
 array([0.          , 0.67142857, 1.          ]),
 array([2, 1, 0], dtype=int64))

```


In [91]:

```

1 # Let's create columns with different probability cutoffs
2 numbers = [float(x)/10 for x in range(10)]
3 for i in numbers:
4     y_pred_final[i]= y_pred_final.Survived_Prob.map( lambda x: 1 if x > i else 0)
5 y_pred_final.head()

```

Out[91]:

	survived	name	Survived_Prob	predicted	0.0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9
0	0.0	173	0.607130	1	1	1	1	1	1	1	1	0	0	0
1	0.0	843	0.075576	0	1	0	0	0	0	0	0	0	0	0
2	0.0	996	0.088081	0	1	0	0	0	0	0	0	0	0	0
3	0.0	992	0.605983	1	1	1	1	1	1	1	1	0	0	0
4	1.0	12	0.963953	1	1	1	1	1	1	1	1	1	1	1

In [92]:

```

1 # Now Let's calculate accuracy sensitivity and specificity for various probability cutoffs
2 cutoff_df = pd.DataFrame( columns = ['prob','accuracy','sensi','speci'])
3 from sklearn.metrics import confusion_matrix
4 num = [0.0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9]
5 for i in num:
6     cm1 = metrics.confusion_matrix( y_pred_final.survived, y_pred_final[i] )
7     total1=sum(sum(cm1))
8     accuracy = (cm1[0,0]+cm1[1,1])/total1
9     sensi = cm1[0,0]/(cm1[0,0]+cm1[0,1])
10    speci = cm1[1,1]/(cm1[1,0]+cm1[1,1])
11    cutoff_df.loc[i] =[ i ,accuracy,sensi,speci]
12 print(cutoff_df)

```


	prob	accuracy	sensi	speci
0.0	0.0	0.356234	0.000000	1.000000
0.1	0.1	0.526718	0.296443	0.942857
0.2	0.2	0.681934	0.592885	0.842857
0.3	0.3	0.732824	0.699605	0.792857
0.4	0.4	0.732824	0.743083	0.714286
0.5	0.5	0.768448	0.822134	0.671429
0.6	0.6	0.788804	0.881423	0.621429
0.7	0.7	0.793893	0.956522	0.500000
0.8	0.8	0.791349	0.972332	0.464286
0.9	0.9	0.717557	0.988142	0.228571

In [93]:

```
1 # Let's plot accuracy sensitivity and specificity for various probabilities.
2 cutoff_df.plot.line(x='prob', y=['accuracy', 'sensi', 'speci'])
```

Out[93]:

<matplotlib.axes._subplots.AxesSubplot at 0x215b1629e88>



In [95]:

```
1 y_pred_final['final_predicted'] = y_pred_final.Survived_Prob.map( lambda x: 1 if x > 0
```

In [96]:

```
1 y_pred_final.head()
```

Out[96]:

	survived	name	Survived_Prob	predicted	0.0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	fir
0	0.0	173	0.607130	1	1	1	1	1	1	1	1	0	0	0	
1	0.0	843	0.075576	0	1	0	0	0	0	0	0	0	0	0	
2	0.0	996	0.088081	0	1	0	0	0	0	0	0	0	0	0	
3	0.0	992	0.605983	1	1	1	1	1	1	1	1	0	0	0	
4	1.0	12	0.963953	1	1	1	1	1	1	1	1	1	1	1	

In [98]:

```
1 #Let's check the overall accuracy.
2 metrics.accuracy_score( y_pred_final.survived, y_pred_final.final_predicted)
```

Out[98]:

0.732824427480916

In [100]:

```
1 metrics.confusion_matrix( y_pred_final.survived, y_pred_final.final_predicted )
```

Out[100]:

```
array([[177,  76],
       [ 29, 111]], dtype=int64)
```

In []:

```
1
```